QueryGenie

A Self-Correcting, Confidence-Aware Natural Language to SQL Interface

Ask a database in plain English — and be told when the answer can't be trusted

STUDENTS | **INNOVATION** | **RESEARCH** | **IMPACT**

Python 3.13 · PyTorch · Hugging Face · SQLite · PostgreSQL · MySQL · macOS · Windows app · 45 tests · Open source · GitHub

Abstract

QueryGenie lets anyone ask a database questions in plain English. Text-to-SQL models answer every question with the same confidence, so a wrong answer looks exactly like a right one. We reproduce CodeS (SIGMOD 2024), an open-source text-to-SQL model, to 77.8% execution accuracy on Spider (paper: 77.9%), and add confidence-aware abstention: the system runs every candidate query the model writes and checks whether they agree. When they don't, it asks the user instead of guessing. It answers 90.4% of questions and is right on 83.3% of them, with a third fewer wrong answers — a statistically significant gain. It runs fully locally in a desktop app.

Objectives

1. Reproduce CodeS-1B on Spider within the ±5–10% tolerance
2. Find which parts of the CodeS pipeline matter (ablation)
3. Make the system say when it is unsure — tested statistically
4. Measure accuracy on realistically worded questions
5. Deliver a desktop app that runs fully locally

Key Innovations & Contributions

1. Confidence-aware abstention
Run all candidate queries, count how many agree; ask "Did you mean…?" when they don't

2. Faithful, audited reproduction
77.8% vs 77.9% with the official evaluator, after auditing our own harness

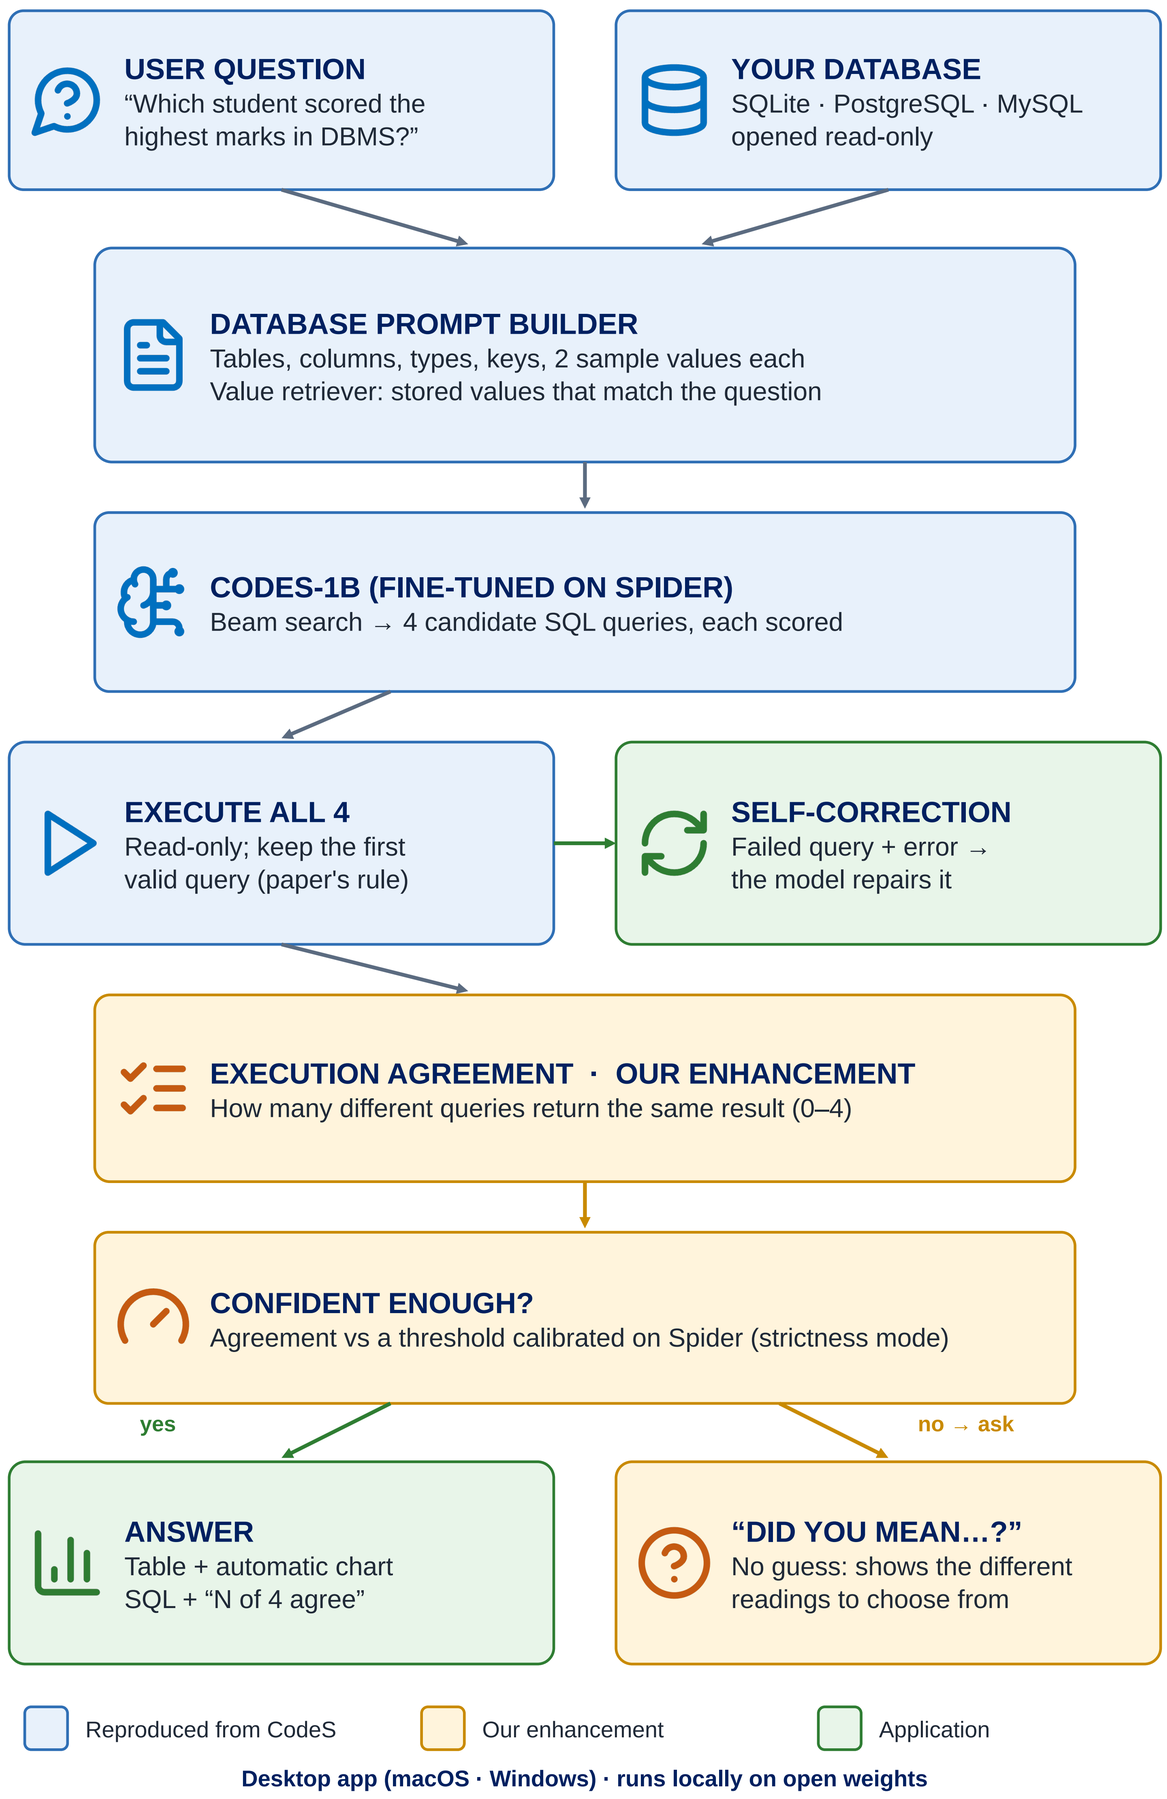
3. Execution-guided self-correction
A failed query goes back to the model with the database error

4. Private by design
Open weights, runs on the user's computer — no schema sent to a third party

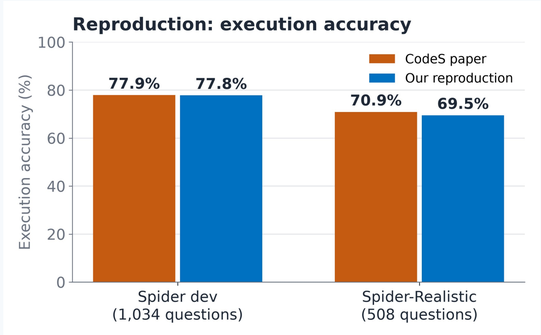
Datasets (Hugging Face — approved source)

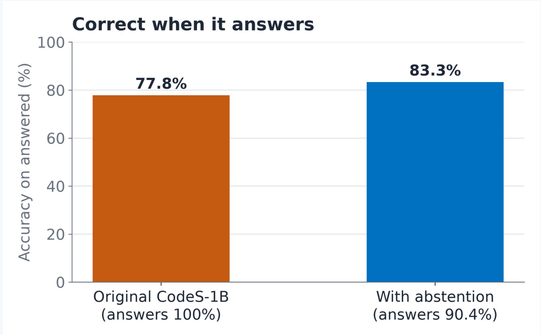
Benchmark	Questions	Used for
Spider dev	1,034 · 20 databases	Reproduction, ablation, enhancement
Spider-Realistic	508	Questions without column names
BIRD Mini-Dev	500 · 11 databases	Real-world databases (future work)

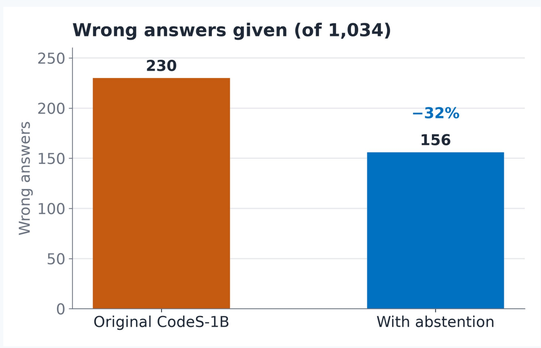
Proposed Methodology / System Architecture

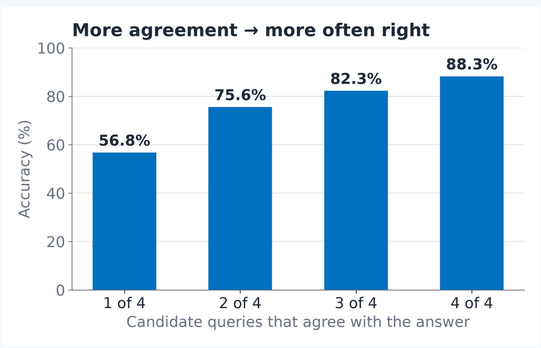


Experimental Results









77.8% Spider EX (paper 77.9%)

69.5% Spider-Realistic (paper 70.9%)

83.3% correct when it answers (vs 77.8%)

−32% wrong answers (230 → 156)

$p < 10^{-9}$ paired t-test

Significant: paired t-test $p = 7.7\times10^{-10}$, Wilcoxon $p = 9.8\times10^{-4}$, bootstrap 95% CI +4.5 to +6.7 points.
Robust: counting each different query only once, accuracy at 90% coverage is still 83.5%.

References

[1] Li, H., et al. (2024). CodeS: Towards building open-source language models for text-to-SQL. arXiv:2402.16347.

[2] Yu, T., et al. (2018). Spider: A large-scale human-labeled dataset for complex and cross-domain semantic parsing and text-to-SQL task. arXiv:1809.08887.

[3] Deng, X., et al. (2020). Structure-grounded pretraining for text-to-SQL. arXiv:2010.12773.

[4] Li, J., et al. (2023). Can LLM already serve as a database interface? A big bench for large-scale database grounded text-to-SQLs. arXiv:2305.03111.

[5] Pourreza, M., & Rafiei, D. (2023). DIN-SQL: Decomposed in-context learning of text-to-SQL with self-correction. arXiv:2304.11015.

[6] Wang, X., et al. (2022). Self-consistency improves chain of thought reasoning in language models. arXiv:2203.11171.

[7] Geifman, Y., & El-Yaniv, R. (2017). Selective classification for deep neural networks. arXiv:1705.08500.

[8] Zhong, R., Yu, T., & Klein, D. (2020). Semantic evaluation for text-to-SQL with distilled test suites. arXiv:2010.02840.

Team Details

Supervisor
M Parameswar
Associate Professor

Team 18
2320090050
N V V S S Jayakanth Kamisetti
2320090009
Boddupally Narendar
2320090018
M Praneeth Reddy

Future Work

- CodeS-3B and more, more diverse candidate queries
- Real-world databases (BIRD)
- A platform for any database (Spark, Snowflake, NoSQL) with access control